We have made the following revisions in response to the reviewers' comments:

- We revised the Introduction to clarify why we perform a Multiword Expression (MWE) identification task on the PARSEME corpus, despite the fact that (i) we present an algorithm for Puns in Multiword Expressions (PMWEs), and (ii) the PARSEME corpus does not explicitly contain PMWEs. Although the corpus lacks PMWEs, we identify several subsets of MWEs that share structural similarities with PMWEs, such as MWE variants, discontinuous MWEs, and unseen-in-train MWEs. This rationale is now clearly explained in the Introduction, and results involving these subsets are analyzed in Section 6.

- We also took the opportunity to clarify other aspects of the Introduction, particularly by disambiguating the concepts of MWEs and PMWEs. While our primary focus is on PMWEs, these are derived from MWEs, and as such, we use MWE-related resources and evaluation methods to assess our algorithm.

- We corrected a minor error in Table 4, where the accuracy value was incorrect for two approaches (fuzzy and combined). Additionally, Table 5 has been reformatted for improved readability by reducing the number of bolded values.

- We added several missing references as requested by the reviewers. We also intend to include additional references once we are permitted to add an extra page of material.

- In response to one reviewer’s request, we included in the Appendix a table with detailed statistics on the PARSEME 1.3 corpus, broken down by language and data partition.

- Addressing the metareviewer’s suggestion that Levenshtein distance should be faster than cosine similarity, we clarify that this is not necessarily the case. Levenshtein distance has a time complexity of O(nm), while cosine similarity has a complexity of O(n). To investigate further, we conducted a small exploratory experiment (unrelated to the main paper) and found both methods to have similar runtime performance. We then re-ran our snowclone detection experiment using Levenshtein distance instead of cosine similarity. The results show that ASMR performs worse with Levenshtein distance, likely because it favors shorter candidates, leading to false positives among the top-ranked results. This additional experiment is detailed in the Appendix.

- We also conducted additional parameter analysis for the snowclone detection task, described in the Appendix. This includes an evaluation of how vectorizer parameters affect performance. We also specify the best parameter configurations for both the snowclone detection and MWE identification tasks.

- As mentioned in our response to the reviewers, we did not incorporate large language models (LLMs) into our study. This decision was based primarily on two reasons: (i) we believe that PMWE identification using LLMs warrants a dedicated study of its own, and (ii) we currently lack the resources to train or fine-tune such models adequately. For these reasons, we chose to focus our efforts on evaluating the ASMR algorithm. In the future, we plan to leverage ASMR to help create the necessary resources for exploring LLM-based approaches.

- One reviewer suggested extending our work by creating a dataset specifically for PMWEs. As noted in our response, we are indeed working on the development of such a dataset. However, building this type of resource is time-consuming and may fall outside the scope of the present paper.

- We added a brief explanation of the alignment methods used, which can be found in the Appendix.

- Finally, we acknowledge a reviewer's observation that our Related Work section lacks a discussion of exact string matching approaches. We agree with this point and plan to include such a paragraph once we are allowed to expand the paper by one additional page.

We would like to thank the reviewers and the metareviewer for their thorough and constructive feedback. Their insightful comments helped us improve the clarity, precision, and overall quality of the paper. We greatly appreciate the time and effort they dedicated to reviewing our work.